


# OpenMP

---

# Shared Memory Architektur

---



# OpenMP

---

- Portable programming of shared memory systems.
- It is a quasi-standard.
- OpenMP-Forum
- 1997 - ...
- API for Fortran and C/C++
  - directives
  - runtime routines
  - environment variables
- [www.openmp.org](http://www.openmp.org)

# Example

---

## Program

```
#include <omp.h>

main() {
    #pragma omp parallel
    {
        printf("Hello world");
    }
}
```

## Compilation

```
> icc -O3 -openmp openmp.c
```


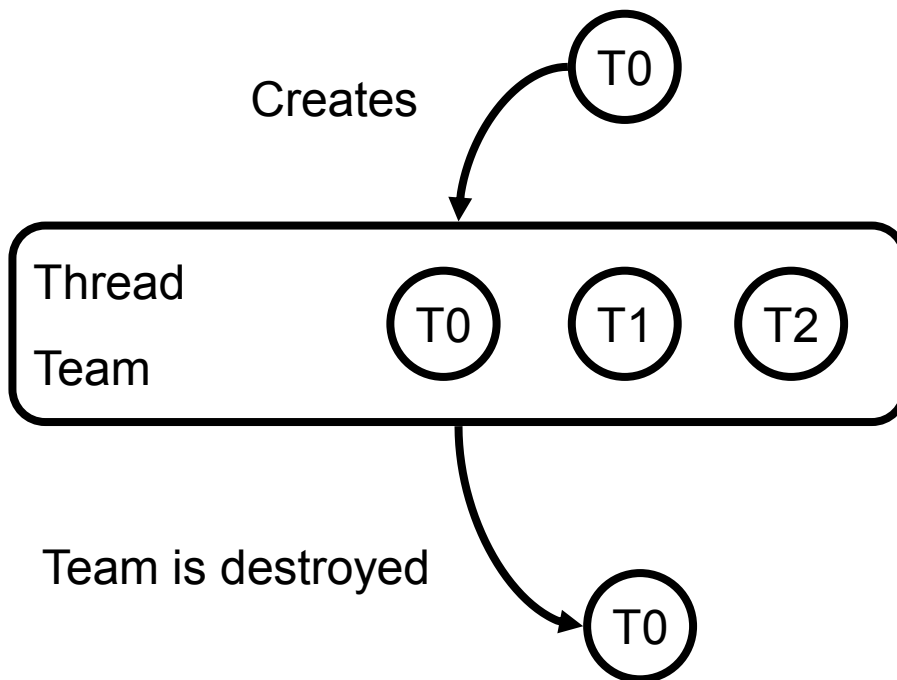
## Execution

```
> export OMP_NUM_THREADS=2
> a.out
Hello world
Hello world
```

```
> export OMP_NUM_THREADS=3
> a.out
Hello World
Hello World
Hello World
```

# Execution Model

```
#pragma omp parallel
{
    printf("Hello world %d\n", omp_get_thread_num());
}
```



# Fork/Join Execution Model

---

1. An OpenMP-program starts as a single thread (*master thread*).
  2. Additional threads (*Team*) are created when the master hits a parallel region.
  3. When all threads finished the parallel region, the new threads are given back to the runtime or operating system.
- A team consists of a fixed set of threads executing the parallel region redundantly.
  - All threads in the team are synchronized at the end of a parallel region via a barrier.
  - The master continues after the parallel region.

# Work Sharing in a Parallel Region

---

```
main () {  
    int a[100];  
    #pragma omp parallel  
    {  
        #pragma omp for  
        for (int i= 1; i<n;i++)  
            a(i) = i;  
        ...  
    }  
}
```

# Shared and Private Data

---

- Shared data are accessible by all threads. A reference `a[5]` to a shared array accesses the same address in all threads.
- Private data are accessible only by a single thread. Each thread has its own copy.
- The default is shared.



# Private clause for parallel loop

---

```
main () {  
    int a[100], t;  
  
    #pragma omp parallel  
    {  
        #pragma omp for private(t)  
        for (int i= 1; i<n;i++){  
            t=f(i) ;  
            a(i)=t;  
        }  
    }  
}
```

# Private Data

---

- A new copy is created for each thread.
- One thread may reuse the global shared copy.
- The private copies are destroyed after the parallel region.
- The value of the shared copy is undefined.

# Parallel Region

---

- The statements enclosed lexically within a region define the **lexical extent** of the region.
- The **dynamic extent** further includes the routines called from within the construct.

```
#pragma omp parallel [parameters]
{
    parallel region
}
```

# Work-Sharing Constructs

---

- Work-sharing constructs distribute the specified work to all threads within the current team.
- Types
  - Parallel loop
  - Parallel section
  - Master region
  - Single region
  - General work-sharing construct (only Fortran)

# Parallel Loop

---

```
#pragma omp for [parameters]  
  for ...
```

- The iterations of the do-loop are distributed to the threads.
- The scheduling of loop iterations is determined by one of the scheduling strategies *static*, *dynamic*, *guided*, and *runtime*.
- There is no synchronization at the beginning.
- All threads of the team synchronize at an implicit barrier if the parameter *nowait* is not specified.
- The loop variable is by default private. It must not be modified in the loop body.
- The expressions in the for-statement are very restricted.

# Scheduling Strategies

---

- **Schedule clause**

`schedule (type [,size])`

- **Scheduling types:**

- **static**: Chunks of the specified size are assigned in a round-robin fashion to the threads.
- **dynamic**: The iterations are broken into chunks of the specified size. When a thread finishes the execution of a chunk, the next chunk is assigned to that thread.
- **guided**: Similar to dynamic, but the size of the chunks is exponentially decreasing. The size parameter specifies the smallest chunk. The initial chunk is implementation dependent.
- **runtime**: The scheduling type and the chunk size is determined via environment variables.

# Example: Dynamic Scheduling

---

```
main() {  
    int i, a[1000];  
  
    #pragma omp parallel  
    {  
        #pragma omp for schedule(dynamic, 4)  
        for (int i=0; i<1000;i++)  
            a[i] = omp_get_thread_num();  
  
        #pragma omp for schedule(guided)  
        for (int i=0; i<1000;i++)  
            a[i] = omp_get_thread_num();  
    }  
}
```

# Reductions

---

```
reduction(operator: list)
```

- This clause performs a reduction on the variables that appear in *list*, with the operator *operator*.
- Variables must be shared scalars
- *operator* is one of the following:
  - +, \*, -, &, ^, |, &&, ||
- Reduction variable might only appear in statements with the following form:
  - $x = x \text{ operator } \text{expr}$
  - $x \text{ binop} = \text{expr}$
  - $x++$ ,  $++x$ ,  $x--$ ,  $--x$



# Example: Reduction

---

```
#pragma omp parallel for reduction(+: a)
for (i=0; i<n; i++) {
    a = a + b[i];
}
```

# Classification of Variables

---

- `private(var-list)`
  - Variables in var-list are private.
- `shared(var-list)`
  - Variables in var-list are shared.

# Parallel Section

---

```
#pragma omp sections [parameters]
{
    [#pragma omp section]
        block
    [#pragma omp section]
        block ]
}
```

- Each section of a parallel section is executed once by one thread of the team.
- Threads that finished their section wait at the implicit barrier at the end of the section construct.

# Example: Parallel Section

---

```
main() {  
    int i, a[1000], b[1000]  
  
    #pragma omp parallel private(i)  
    {  
        #pragma omp sections  
        {  
            #pragma omp section  
            for (int i=0; i<1000; i++)  
                a[i] = 100;  
            #pragma omp section  
            for (int i=0; i<1000; i++)  
                b[i] = 200;  
        }  
    }  
}
```

# Master / Single Region

---

```
#pragma omp master  
    block  
  
#pragma omp single [parameters]  
    block
```

- A master or single region enforces that only a single thread executes the enclosed code within a parallel region.
- Common
  - No synchronization at the beginning of region.
- Different
  - Master region is executed by master thread while the single region can be executed by any thread.
  - Master region is skipped by other threads while all threads are synchronized at the end of a single region.

# Combined Work-Sharing and Parallel Constructs

---

- `#pragma omp parallel for`
- `#pragma omp parallel sections`
- `!$OMP PARALLEL WORKSHARE`

# Barrier

---

```
#pragma omp barrier
```

- The barrier synchronizes all the threads in a team.
- When encountered, each thread waits until all of the other threads in that team have reached this point.

# Critical Section

---

```
#pragma omp critical [(Name)]  
{ ... }
```

- **Mutual exclusion**

- A critical section is a block of code that can be executed by only one thread at a time.

- **Critical section name**

- A thread waits at the beginning of a critical section until no other thread is executing a critical section with the same name.
- All unnamed critical directives map to the same name.
- Critical section names are global entities of the program. If a name conflicts with any other entity, the behavior of the program is unspecified.



# Simple Locks

---

- Locks can be hold by only one thread at a time.
- A lock is represented by a lock variable of type `omp_lock_t`.
- The thread that obtained a simple lock cannot set it again.
- Operations
  - `omp_init_lock(&lockvar)`: initialize a lock
  - `omp_destroy_lock(&lockvar)`: destroy a lock
  - `omp_set_lock(&lockvar)`: set lock
  - `omp_unset_lock(&lockvar)`: free lock
  - `logicalvar = omp_test_lock(&lockvar)`: check lock and possibly set lock, returns true if lock was set by the executing thread.

# Explicit Tasking

---

- Explicit creation of tasks

```
#pragma omp parallel
{
    #pragma omp single {
        for ( elem = l->first; elem; elem = elem->next)
            #pragma omp task
                process(elem)
    }
    // all tasks are complete by this point
}
```

- Task scheduling

- Tasks can be executed by any thread in the team

- Barrier

- All tasks created in the parallel region have to be finished.

# Summary

---

- OpenMP is quasi-standard for shared memory programming
- Based on Fork-Join Model
- Parallel region and work sharing constructs
  - Declaration of private or shared variables
  - Reduction variables
  - Scheduling strategies
- Synchronization via Barrier, Critical section, Atomic, locks, nestable locks