

Lab Course: Efficient Programming of Multicore Processors and Supercomputers

Group 1: Jonas Mayer, Paul Preißner, Konrad Pröll

July 18, 2017

6.3 Bad Parallelization of AlphaBeta

1. Comparison of total number of evaluations of minimax and abid

Setting	MiniMax	AlphaBeta
midgame1, depth2	5,199	4,866
midgame1, depth3	120,718	49,198
midgame1, depth4	2,485,542	916,465
midgame1, depth5	15,018,100	4,120,198
midgame2, depth2	4,358	4,577
midgame2, depth3	97,908	43,335
midgame2, depth4	1,519,234	531,259
midgame2, depth5	4,411,737	2,067,878
endgame, depth2	1,796	3,602
endgame, depth3	17,823	9,801
endgame, depth4	23,266	17,643
endgame, depth5	1,520,313	402,920

There are two things one can easily spot: For small depths, alphabeta might actually need more evaluations than minimax. This assumption, however, is not correct, as a pure alphabeta cannot evaluate more moves than minimax. The overhead in this case is caused by optimizations applied to alphabeta, mainly iterative deepening, but also by first trying to initialize the alphabeta window with reasonable values, as there is not a lot to optimize for such flat trees. Second, you can clearly spot, that those optimizations as well as the general alphabeta approach do a good job when it comes to reducing the size of search trees with huge depth.

2. Naive Parallelization, Overhead

We decided to decompose the moves at the topmost layer in the alphabeta method and leave the searchbestmove method untouched, as it does things we

cannot really parallelize (to parallelize iterative deepening is clearly not advantageous, because the idea is to reuse the results from the calculation for depth-1 when calculating the values for depth, and parallelizing the alphabeta-calls with different windows is not smart, because we assume that the reasonable values are actually reasonable, and you can see from the fact that there is not a lot of calls of alphabeta with the same maxdepth that they are.) We used a master/slave(worker) approach, so that the root process distributes one move at a time to a process, and also sent the corresponding alphabeta window to each process with the particular move. We did this, because the window might have updated and we want to allow processes to prune quickly if they can.

Setting	AlphaBeta	Overhead 2 workers	4 processes (3w)	6p (5w)	8p (7w)
midgame1, depth5	4,120,198	18.8%	19.5%	24.8%	22.6%
midgame2, depth5	2,067,878	6.6%	29.6%	44.7%	59.2%
endgame, depth5	402,920	162%	300.1%	631%	644.5%

The reason for the search overhead is easy to understand. When the sequential version works through the tree, it already knows the value of node1 at any depth while evaluating node2 (assuming that node2 is on the righthand side of node1, i.e. we are enumerating from left to right) at the very same depth and therefore can easily decide whether to prune or not. In a parallel version, there are n nodes (with n being the number of workers) which are being evaluated with the same alphabeta values, because they are running concurrently. Those processes therefore have older values to use when deciding whether to prune or not, and have to do more evaluations. In particular cases, where predicting the alphabeta window while using iterative deepening fails, you will have a large amount of overhead as there are going to be depthsearches with a broad window and high depth. Where a sequential algorithm only has to traverse the first node with the broad window, the parallel version must traverse n (with n again being the number of workers) nodes with the broad window (or, in our case n+1, as the masterprocess sends out a new task before reading the values returned by the worker), resulting in plenty of overhead. Exactly this issue happens in position-endgame, leading to the incredible amount of overhead. The values are only computed for more than one worker, but due to the master/worker approach, you could actually expect overhead for one worker as well, because the master process already sends out a new task before processing the data returned by the worker.

Speedup

In general, the result is better than expected, mainly because the iterative deepening and adaptive alphabeta-window do a good job narrowing down search windows for depthsearches with high levels. However, you can see the devastating effects the mispredicted alphabeta-window for position-endgame has.

Speedup for ab using naive approach at depth 5

